# Implementing Map-Scan Fusion in the Futhark Compiler

Bachelor project

Brian Spiegelhauer
brianspieg@gmail.com

William Jack Lysgaard Sprent
bsprent@gmail.com

June 2, 2016

Contents

1 Abstract

2 Introduction

NOTE: the contents of this section is lifted from our synopsis and is probably placeholder
The Futhark language is a functional programming with which the main idea is to allow for the expression of sufficiently complex programs while keeping complexity to a level where programs can be aggressively optimised and have their parallelism exploited [1].

The Futhark compiler already supports a range of fusion optimisations [2], but does not currently support fusion between `Map` and `Scan` statements.

For our project we will explore the possibility of implementing Map-Scan fusion into the Futhark compiler, and will examine the performance benefits (if any) of performing such optimisations.

2.1 Motivation

Fusion has the "[..] potential to optimize both the memory hierarchy time overhead and, sometimes asymptotically, the space requirement" [2]. Hence the main motivation for adding Map-Scan fusion capabilities to the optimiser of the Futhark compiler, is the potential for enabling performance increases for some Futhark programs.

2.2 Tasks

The project can be divided into three main tasks:

1. Gain an understanding of logical reasoning behind fusion optimisations on Second Order Array Combinators.

2. Read and understand the relevant parts of the Futhark compiler required to make the necessary changes in the compiler.

3. Modify all modules of the Futhark compiler necessary to implement the Map-Scan fusion itself.

At first sight, these tasks look fairly straight forward. However, we expect that the main difficulties of this project lie within unforeseen roadblocks we will run into when modifying the codebase.

3 Background Information

There are problems/calculations that gets to a size where normal sequential programming involving consecutive execution of processes, will reach a computation time unsatisfying for the intended users. In some cases these calculations can be done much faster with parallel programming. Parallel programming is where many calculations are carried out simultaneously, with the idea of dividing a problem into smaller sub problems solved at the same time. Parallel and sequential programming are not mutually exclusive, in the sense that if you use parallel programming, you cant use sequential programming, in many cases they are used together. Parallel programming can be done on the CPU with its multiple cores, but when possible and advantages it is much better to harness the thousands of cores in the GPU - graphics processing units. The GPU is no longer only used to do graphical calculations, but also General-purpose computing, GPGPU (General-purpose computing on graphics processing units).

To do GPGPU, Hyperfit a joint research center addressing the simultaneous challenges of high transparency, high computational performance and high productivity in finance, employing an integrated approach of financial mathematics, domain specific languages, parallel functional programming, and high-performance systems [3] created Futhark.

3.1 Futhark

As described in Troels Henriksens master thesis [4, The $\mathcal{L}_0$ language, p. 8] the language $\mathcal{L}_0$ later renamed Futhark is in a sense "sufficient", in that it is Turing-complete, and can express imperative style loops with do-loops. However Futhark is ment to use second-order array combinations (SOACs) to do bulk operations on arrays instead of using the do-loops. In this sections the reasoning behind using SOACs will be explained by showing the difference in their computation when done sequentially vs. parallelly.

3.2 SOACs

Both `map` and `scan` are defined as SOACs – or Second Order Array Combinators. Hence they have no free variables, take first-order functions as arguments, and output first-order functions whose domains are arrays of the domain of the input. Furthermore, in Futhark, these array inputs and outputs are tuples of arrays, and not arrays of tuples. While Futhark SOACs are semantically indentical to their namesake higer-order functions found in other functional languages, working with SOACs allows for some assumptions to be made which turn out to be useful in regards to both parallisation and optimisation. In particular each SOAC can be considered as representing a specific shape of an imperative do-loop, which is used in Futhark to expedite the loop-fusion process. [4, chap. 7]

3.2.1 Map

The `map` $f\ a$ function, has the very simple definition of taking a function $f\ :\ \alpha \rightarrow \beta$ and returning a function `map` $f\ :\ [\alpha] \rightarrow [\beta]$ which applies f to every element of an input array, a. This gives us the type signature of `map`,

$$\texttt{map}\ f\ a\ :\ (\alpha \rightarrow \beta) \rightarrow [\alpha] \rightarrow [\beta].$$

bAnd the semantic definition of `map`,

$$\texttt{map}\ f\ a\ =\ [f(a_0), f(a_1), ..., f(a_{n-1})].$$

Having no free variables, means that each result $f(a_i)$ *only* depends on the corresponding element a_i. This makes `map`s fantastic for parallelisation as once the degree of parallism reaches the size of a, `map` $f\ a$ can be potentially be computed in a single parrallel step, or c steps for a chunk size of c.

3.2.2 Scan

$\texttt{scan} \odot e\, a$ takes a binary, associative function $\odot\, :\, \alpha \to \alpha \to \alpha$ and returns a function $\texttt{scan} \odot\, :\, \alpha \to [\alpha] \to [\alpha]$ which computes the $\odot$ prefixes of an input array a starting with a neutral element, e. Overall, $\texttt{scan}$ has the type signature,

$$\texttt{scan} \odot e\, a\, :\, (\alpha \to \alpha \to \alpha) \to \alpha \to [\alpha] \to [\alpha].$$

Computing $\texttt{scan}$ with the function $\odot$, the array a, and neutral element e gives us,

$$\texttt{scan} \odot e\, a\, = [e \odot a_0, e \odot a_0 \odot a_1, ..., e \odot a_0 \odot ... \odot a_{n-1}].$$

However, computing such a $\texttt{scan}$ is not as simple as with a $\texttt{map}$ as each prefix $a_0 \odot ... \odot a_i$ obviously depends on the previous prefix $a_0 \odot ... \odot a_{i-1}$. Hence, the associativity of $\odot$ is vital as it means that this dependency does not force computation order, and partial results can be computed independently and combined.

4 Map-Scan Fusion

In a typical situation involving $\texttt{map}$ and $\texttt{scan}$, we will have the $\texttt{map}$ producing some array, which is subsequently consumed by the $\texttt{scan}$:

$$b = \texttt{map}\, f\, a$$
$$c = \texttt{scan} \odot e\, b.$$

In this situation, the memory access pattern will look something like the following,

$$\texttt{load}[a_0], \texttt{store}[b_0], \texttt{load}[a_1], \texttt{store}[b_1], ..., \texttt{load}[a_{n-1}], \texttt{store}[b_{n-1}]$$
$$\texttt{load}[b_0], \texttt{store}[c_0], \texttt{load}[b_1], \texttt{store}[c_1], ..., \texttt{load}[b_{n-1}], \texttt{store}[c_{n-1}].$$

Recalling that element of the scanned array c depend only on the previous prefixes, we have that c_i will only depend on b_j for $j \leq i$ being calculated. This means that, ideally, we can interleave this access pattern without breaking any dependencies – giving us a pattern more like this:

$$\texttt{load}[a_0], \texttt{store}[b_0], \texttt{load}[b_0], \texttt{store}[c_0], ..., \texttt{load}[a_{n-1}], \texttt{store}[b_{n-1}], \texttt{load}[b_{n-1}], \texttt{store}[c_{n-1}].$$

In this case, all accesses to b become wholly uneccessary and can be removed, giving us:

$$\texttt{load}[a_0], \texttt{store}[c_0], ..., \texttt{load}[a_{n-1}], \texttt{store}[c_{n-1}].$$

This is the goal of perfoming $\texttt{map-scan}$ fusion. By fusing a $\texttt{map}$ into a $\texttt{scan}$ we can optimise a program by eliminating costly memory accesses. Something which is even more useful when working on a GPGPU, which generally do not have the more forgiving memory hierarchy and caching system of a CPU.

4.1 Scanomap

When looking to fuse a $\texttt{map}$ into a $\texttt{scan}$, the most straight forward approach is to attempt to perform function composition on the input functions of the respective functions. Hence, turning the following

$$b = \texttt{map}\, f\, a \tag{1}$$
$$c = \texttt{scan} \odot e\, b \tag{2}$$

into,

$$c = \texttt{scan } \odot_f \ e \ a \tag{3}$$

where,

$$x \odot_f y = \odot \circ f = x \odot f(y)$$

would be the naive approach. However, the type signature of the resulting function

$$\odot_f \ : \ \alpha \to \beta \to \alpha$$

is not compatible with the Futhark definition of $\texttt{scan}$, and neither is it associative. Clearly, a different approach is needed.

The solution to this problem is to exploit how Futhark uses chunking. Since the associativity of the $\texttt{scan}$ operator only comes into play during the parallel phase of computation, we can distiguish between how each chunk is computed sequentially and how chunks are joined.

To do so, we must expand the list of SOACs with the internal $\texttt{scanomap}$ function. $\texttt{scanomap}$ is semantically similar to $\texttt{scan}$, however it takes two function parameters – an associatve scanning function meant for parallelly scanning across chunks, and a sequential folding function meant for scanning within a chunk. Scanomap has the following type signature,

$$\texttt{scanomap } \odot \ \odot_f \, e \, a \ : \ (\alpha \to \alpha \to \alpha) \to (\alpha \to \beta \to \alpha) \to \alpha \to [\beta] \to [\alpha].$$

and can be considered semantically similar to performing a left fold with the $\odot_f$ function

$$\texttt{scanomap } \odot \ \odot_f \, e \, a \ = [e \odot_f a_0, (e \odot_f a_0) \odot_f a_1, ..., ((e \odot_f a_0) \odot_f ...) \odot_f a_{n-1}]$$

which also corresponds to the chuck-wise computation of $\texttt{scanomap}$. The scanning operator can then be used to join two chunks,

$$\texttt{scanomap } \odot \ \odot_f \, e \, (a +\!\!+ \, b) \ = [a'_0, a'_1, ..., a'_{n-1}] +\!\!+ \ [b'_0 \odot a'_{n-1}, b'_1 \odot a'_{n-1}, ..., b'_{n-1} \odot a'_{n-1}]$$

where

$$\texttt{scanomap } \odot \ \odot_f \, e \, a \ = [a'_0, a'_1, ..., a'_{n-1}]$$
$$\texttt{scanomap } \odot \ \odot_f \, e \, b \ = [b'_0, b'_1, ..., b'_{n-1}].$$

For $\texttt{scanomap}$ to be used in facilitating $\texttt{map-scan}$ fusion, we first observe that we have the following equivalence,

$$\texttt{scan } \odot \ e \, a \equiv \texttt{scanomap } \odot \ \odot \, e \, a.$$

Hence we can freely turn any regular $\texttt{scan}$ into an equivalent $\texttt{scanomap}$. If we once again take a look at our previous example we can see how we can now fuse a $\texttt{map}$ into a $\texttt{scan}$ using the $\texttt{scanomap}$ construction:

$$b = \texttt{map } f \, a$$
$$c = \texttt{scan } \odot \ e \, b$$

Firstly, we turn the $\texttt{scan}$ into an equivalent $\texttt{scanomap}$,

$$c = \texttt{scanomap } \odot \ \odot \, e \, b$$

and then compose a folding function, $\odot \circ f = \odot_f$, from the mapping function and the scanning operator and discard the intermediate b list, to arrive at an equivalent $\texttt{scanomap}$,

$$c = \texttt{scanomap } \odot \ \odot_f \, e \, a.$$

Any producer $\texttt{map}$ can be fused into a consuming $\texttt{scan}$ in this manner.

4.2   Necessary Conditions

There are a series of conditions for the fusion transformation can happen. As well as a number of cases where a fusion should not happen. The fusion transformation assumes a normalized program, with the following properties: [2, Figure 5, page 4]

- No tuple type can appear in an array or tuple type, i.e., flat tuples,

- unzip has been eliminated, zip has been replaced with assertZip, which verifies either statically or at runtime that the outer size of zip's input matches, and finally, the original SOACs (map) have been replaced with their tuple-of-array version,

- tuple expressions can appear only as the final result of a function, SOAC, or if expression, and similarly for the tuple pattern of a let binding, e.g., a formal argument cannot be a tuple,

- e_1 cannot be a let expression when used in let $p = e_1$ in e_2,

- each if is bound to a corresponding let expression, and an if's condition cannot be in itself an if expression, e.g.,
  $a + $ if(if c_1 then e_1 else e_2) then e_3 else $e_4 \rightarrow$
  let $c_2 = $ if c_1 then e_1 else e_2 in
  let $b = $ if c_2 then e_3 else e_4 in $a + b$

- function calls, including SOACs, have their own let binding, e.g.,
  reduce2(f, a) $+ x \Rightarrow$ let $y = $ reduce2(f, e, a) in $y + x$,

- all actual arguments are vars, e.g., $f(a + b)) \Rightarrow$ let $x = a + b$ in $f(x)$.

The properties listed above is reached by going through the different stages in the compiler pipeline, as shown in figure 1. Type checking first so that any errors will refer to the actual names given by the programmer, then moving on to tuple transformation that flattens all tuples and converts arrays of tuples to tuples of arrays. Tuple and let normalisation to optimise the code based on code recognition (e.g $f(a + b)) \Rightarrow$ let $x = a + b$ in $f(x))$.


Figure 1: Compile pipeline [2, page 4]

When the program has reached the properties mentioned it enters the enabling optimisations loop, with aggressive in lining (i.e., building the call-graph and moving non recursive functions into the function calling it.). Copy propagation e.g.,

$$y = x$$

$$z = 5 * y$$

Would become

$$z = 5 * x$$

and constant folding where the compiler evaluating constant expressions at compile time. After dead code removal the process repeats unless a fixed point is reached, if the fixed point is reached the program is ready for possible fusion transformations. Reaching the fusion stage does not guarantee that a fusion transformation can be done, as shown in figure 2 there are 6 cases where a fusion transformation should not happen.

```
// Case 1: don't fuse if it      // Case 2: not all SOAC
//moves an array use across      //combinations are fusable
//its in-place update point      let x = filter2(c₁, a) in
let x   = map2(f, a) in          let y = filter2(c₁, b) in
let a[1]= 3.33         in        let z = map2   (f,x,y) in..
let y   = map2(g, x) in ..
                                 // Case 4: don't fuse in 2
// Case 3: don't fuse from       //kernels sharing a CF path
//outside in a loop (or λ)       let x = map2(f, arr)  in
let x = map2(f, a)   in          let y = if c then map2(g₁,x)
loop(arr) = for i < N do                      else map2(g₂,x)
  map2(op +, arr, x)             in let z = map2(h, x) in ..

// Case 5: don't fuse if x       // Case 6: x & y used exactly
//used outside SOAC inputs       //once but still don't fuse
let x = map2(f, a) in            let (x,y) = map2(f, a)   in
let y = map2(g, x)   in          let u = reduce2(op +,0,x)in
    x[i] + y[i]                  let v = reduce2(op *,1,y)..
```

Figure 2: Don't Fuse Cases [2, page 6]

Case 1: Fusion across a in-place update is not possible. When fusing the producer array is not created as before, but a part of the computation in the resulting fused SOAC.

Case 2: Not all combinations of SOACs are fusable.

Case 3: Fusion across a loop or a SOAC lambda would duplicate the computation and potentially change the time complexity of the program. In this case instead of calculating variable x once, a fusion would result in x being calculated in each of the loops.

Case 4: When the array x produced by a SOAC is consumed by two other SOACs located on the same execution path, the computation will be duplicated.

Case 5: If array x produced by SOAC is used other then input to another SOAC fusion is not allowed.

Case 6: If two arrays created by a SOAC is used as input in more then one SOAC the computaion is duplicated, and fusion is therefore not allowed.

However there are special cases for scanomap, as the result array x of the producer SOAC in the scanomap transformation, also will be returned by the scanomap in cases where other SOACs consumes array x.

4.3 Fusing Scanomap

Map-Scanomap Fusion   Fusing a `map` as a producer into a consumer `scanomap` is a simple process. Given a `map` with function g and a `scanomap` with the folding function $\odot_f$, we compose a new folding function $x \odot_{f \circ g} y = x \odot f(g(y))$. This is illustrated below,

$$b = \mathtt{map}\, g\, a$$
$$c = \mathtt{scanomap}\, \odot\, \odot_f\, e\, b$$
$$\Downarrow$$
$$c = \mathtt{scanomap}\, \odot\, \odot_{f \circ g}\, e\, a$$

Scanomap-Scanomap Fusion   Horizontally fusing a `scanomap` with another `scanomap`, requires that theres exists no dependency between the two SOACs – i.e. there exists no producer-consumer relationship between the two. In theses cases, it may still prove beneficial to perform fusion – even though there is no direct optimisation – to enable more fusion. This

kind of fusion is performed by in a way concatenating input, output, and both functions for the input `scanomaps`, which gives the following,

$$c = \texttt{scanomap} \odot_1 \odot_{1f} e_1\, a$$
$$d = \texttt{scanomap} \odot_2 \odot_{2g} e_2\, b$$
$$\Downarrow$$
$$(c, d) = \texttt{scanomap} \odot' \odot'_{fg} (e_1, e_2)\, (a, b).$$

Where $\odot'(e_1, e_2, x, y) = (e_1 \odot x, e_2 \odot y)$ and $\odot'_{fg}(e_1, e_2, x, y) = (e_1 \odot f(x), e_2 \odot f(y))$. In this way, the resulting `scanomap` now describes the computations of both input `scanomap`s.

5 Implementation

This section will detail the specifics of performing `map-scan`, and other `scanomap` fusions, as well as detailing the process of analysing a program for valid fusions.

5.1 Bottom Up Analysis

Futhark's fusion module gathers fusible SOACs by using a bottom-up analysis pass of the program's AST. This is done by traversing the AST depth first, while using a set of environments and data structures, to among other things keep track created kernels and the production and consumption of arrays. [2]

These are structures and environments are filled out during the depth first traversal of the AST, such that when analysing a specfic SOAC in the program it is known which SOACs consume its outputs, and whether in-place update in between and so on.

For example, when analysing the program found in Figure 3 and the looking at the `map` statement, two simple look-ups will show that while b is consumed by the `scan`, it is consumed by an in-place updated previously. Hence, fusion is not possible.

$$\vdots$$
$$\textbf{let } b = \texttt{map } f\, a \textbf{ in}$$
$$\textbf{let } b[1] = 4 \textbf{ in}$$
$$\textbf{let } c = \texttt{scan} \odot 0\, b \textbf{ in}$$
$$\vdots$$

Figure 3: Infusible program snippet.

Fusion is then attempted along the way whenever the fusion module meets a new SOACs during its analysis. Consequently, we have that fusion is also performed bottom-up, such that the compiler will prioritize potential fusion by how "deep" the candidates occur in the program.

5.2 Map-Scanomap Fusion

Having identified a single `map` producer and a `scanomap` consumer to perform fusion on, the fusion itself may seem fairly simple. However, this is not entirely the case in most situations, as especially function composition isn't as straight forward as it can seem. There are also two factors which complicate this process:

10

- As previously mentioned, all SOACs in Futhark work with tuples of arrays. This means that a $b = \mathtt{map}\ f\ a$ is actually short hand for $(b0, b1, .., bn) = \mathtt{map}\ f\ (a0, a1, ..., an)$, and not all members of the same tuples are consumed, or produced, in the same place.

- Certain SOACs in Futhark support map-out arrays – i.e. array-outputs from a fused $\mathtt{map}$ which are still output after fusion. This is also supported by $\mathtt{scanomap}$.

This section will go through the algorithm for fusing a single $\mathtt{map}$ into a $\mathtt{scanomap}$ – and by extension the algorithm for fusing a $\mathtt{map}$ into a $\mathtt{scan}$.

The function inputs of SOACs in Futhark consist of a set of parameters, a set of $\mathtt{let}$ bindings, and a body which defines the function output. When looking to create a new folding function for the resulting $\mathtt{scanomap}$ construct, composed entirely from the functions found in the to-be-fused $\mathtt{map}$ and $\mathtt{scanomap}$. Firstly, we will find the new parameter set for

$$f(a0, a1, a2, a3) = \tag{1}$$
$$\quad \textbf{let}\ (x1, x2) = g(a0, a1)\ \textbf{in} \tag{2}$$
$$\quad \textbf{let}\ y1 = a2 + a3\ \textbf{in} \tag{3}$$
$$(x1, x2, y1) \tag{4}$$

Figure 4: Example of a SOAC lambda function.

the new folding function. Given a program, wherein a $\mathtt{map}$ is to be fused with a $\mathtt{scanomap}$ into a new $\mathtt{scanomap}$:

$$b = \mathtt{map}\ f\ a$$
$$d = \mathtt{scanomap}\ \odot\ \odot_g\ ne\ c.$$

Where a, b, c, d are all tuples of arrays, and ne is a tuple of neutral elements.

Fusing these two SOACs into a $\mathtt{scanomap}$ becomes a three step process of

1. determining the parameter list of the output $\mathtt{scanomap}$,

2. determining map-out arrays, and

3. composing a new folding function.

Afterwards a new $\mathtt{scanomap}$ can then be constructed which can replace the input $\mathtt{map}$ and $\mathtt{scanomap}$.

Parameters We have that $\mathtt{map}$ and $\mathtt{scanomap}$ are to be fused, so we must also have $(b0, b1, ..., bn) \cap (c0, c1, ..., cn) \neq \emptyset$. However we do not neccessarily have $(b0, b1, ..., bn) = (c0, c1, ..., cn)$. The parameters for our fused product must then become all of the array input parameters from the $\mathtt{map}$ as well as the array input parameters from the $\mathtt{scanomap}$ which aren't produced by the $\mathtt{map}$. Hence, our new list of parameters become,

$$\text{New-Params} = (a \cup c)/b.$$

Map-Out Arrays Next is to find any map-out arrays which need to be returned. These can potentially consist of some subset of the output arrays of both the $\mathtt{map}$ and the $\mathtt{scanomap}$. The map-out arrays already carried by the input $\mathtt{scanomap}$ are by convention placed last in the list of outputs, so they correspond to the $\mathtt{len}(d) - \mathtt{len}(ne)$ last members of d.

$$\text{map-out}_{scanomap} = \mathtt{drop}(\mathtt{len}(ne), d)$$

The map-out arrays from the input `map` are found by using the fusion environment – which keeps track of where arrays are consumed – through the `unfus` function which filters out any arrays which are not consumed later in the program.

$$\text{map-out}_{map} = \texttt{unfus}(b)$$

Combining these two tuples will then give us the full set of map-out arrays for the fused product,

$$\text{map-out} = m_{out} = \texttt{drop}(\texttt{len}(ne), d) \mathbin{+\!\!+} \texttt{unfus}(b).$$

Function Composition　Composing the folding function for the resulting `scanomap` from the input `map`'s f function, and the input `scanomap`'s $\odot_g$ function is a question of combining their bindings such that the result of the f function is computed first and bound such that it can be used to compute $\odot_g$.

The input functions will both consist of some amount of let bindings followed by a body which consists solely of the variables to be returned. Figure 5 shows generalized functions for both inputs where a_i is the tuple consisting of the ith elements of the arrays contained in the a tuple.

$$
\begin{aligned}
f(a_i) = \qquad\qquad &\qquad\qquad \odot_f (ne, c_i) = \\
\textbf{let } x_1 = e_1 \textbf{ in} &\qquad\qquad \textbf{let } y_1 = e_1 \textbf{ in} \\
\textbf{let } x_2 = e_2 \textbf{ in} &\qquad\qquad \textbf{let } y_2 = e_2 \textbf{ in} \\
\vdots\qquad\quad &\qquad\qquad\qquad \vdots \\
\textbf{let } x_n = e_n \textbf{ in} &\qquad\qquad \textbf{let } y_n = e_n \textbf{ in} \\
b_i\qquad\qquad\quad &\qquad\qquad\quad d_i
\end{aligned}
$$

Figure 5: The SOAC functions f and $\odot_g$ from the input `map` and `scanomap`.

Combining these two functions into a single new folding function $\odot_g \circ f$ (or $\odot_{g \circ f}$) is then a question of first applying all of the bindings from the f function, creating a new binding which binds the body of f to the corresponding parameters of $\odot_g$, and then on to that appending the bindings of $\odot_g$.

The body of $\odot_{f \circ g}$ then becomes the outputs of $\odot_g$ minus any map-out arrays appended with the new map-out arrays as found above.

We can then replace the original `map` and `scanomap` with a new `scanomap`,

$$(d / m_{out}) \mathbin{+\!\!+} m_{out} = \texttt{scanomap} \odot \odot_{f \circ g} \, ne \, ((a \cup c)/b).$$

5.3　Scanomap-Scanomap Fusion

Having two `scanomap` in a futhark program with no dependencies between the two SOACs, and equal length input, the program computation time can be optimised using horizontal fusion.

This section will go through the algorithm for fusing two `scanomap` into one `scanomap`. The fusion algorithm have two steps:

1. Composing a new assosiative operation

2. Composing a new folding function.

$$\odot_g \circ f(ne, ((a \cup c)/b)_i) =$$
$$\textbf{let } x_1 = e_1 \textbf{ in}$$
$$\textbf{let } x_2 = e_2 \textbf{ in}$$
$$\vdots$$
$$\textbf{let } x_n = e_n \textbf{ in}$$
$$\textbf{let } c_i = b_i \textbf{ in}$$
$$\textbf{let } y_1 = e_1 \textbf{ in}$$
$$\textbf{let } y_2 = e_2 \textbf{ in}$$
$$\vdots$$
$$\textbf{let } y_n = e_n \textbf{ in}$$
$$(d_i/m_{out_i}) +\!\!\!+ \ m_{out_i}$$

Figure 6: Resulting function $\odot_{f \circ g}$.

$$c = \texttt{scanomap } \odot_1 \ \odot_{1f} \ e_1 \ a$$
$$d = \texttt{scanomap } \odot_2 \ \odot_{2g} \ e_2 \ b$$
$$\Downarrow$$
$$(c, d) = \texttt{scanomap } \odot' \ \odot'_{fg} \ (e_1, e_2) \ (a, b).$$

Parameters With two `scanomap` to be fused, the goal is to join the computations within the same loop. Hence a does not have to equal b, but the the length of the input arrays must be the same. Our new list of parameters simply become:

$$\text{New-Params} = (a, b)$$

New folding function and associative operation As the two `scanomap` are independent of each other the composition of the new fused folding function is fairly easy. Figure 7 show the pre-fusion folding functions for both scanomap, running sequentially after each other. Fusing them is simply adding the additional parameters and adding the function body of one of them into the other as shown in figure 8.

$$\odot_f (ne_1, a_i) = \qquad\qquad \odot_g (ne_2, b_i) =$$
$$\textbf{let } x_1 = e_1 \textbf{ in} \qquad\qquad \textbf{let } y_1 = e_1 \textbf{ in}$$
$$\textbf{let } x_2 = e_2 \textbf{ in} \qquad\qquad \textbf{let } y_2 = e_2 \textbf{ in}$$
$$\vdots \qquad\qquad\qquad\qquad \vdots$$
$$\textbf{let } x_n = e_n \textbf{ in} \qquad\qquad \textbf{let } y_n = e_n \textbf{ in}$$
$$c_i \qquad\qquad\qquad\qquad d_i$$

Figure 7: The SOAC functions $\odot_f$ and $\odot_g$ from the two scanomap

$$\odot'_{fg}\,(ne1, ne2, a_i, b_i) =$$
$$\textbf{let }\ x_1 = e_1\ \textbf{in}$$
$$\textbf{let }\ x_2 = e_2\ \textbf{in}$$
$$\vdots$$
$$\textbf{let }\ x_n = e_n\ \textbf{in}$$
$$\textbf{let }\ y_1 = e_1\ \textbf{in}$$
$$\textbf{let }\ y_2 = e_2\ \textbf{in}$$
$$\vdots$$
$$\textbf{let }\ y_n = e_n\ \textbf{in}$$
$$(c_i, d_i)$$

Figure 8: Fused function $\odot'_{fg}$

In the same manor the assosiative operations is fused into one, enabling the computation of two scanomaps to be done within one loop.

5.4 Other Fusion

6 Benchmarking and Testing

How have we tested our implementation. Does it work? Why?
How does the performance of a fused program compare with a non-fused program both sequentially and parallelly. Why?

7 Conclusion

References

[1] Futhark language documation. `http://futhark.readthedocs.org/en/latest/index.html`. Accessed: 02-03-2016.

[2] Troels Henriksen and Cosmin Eugen Oancea. A t2 graph-reduction approach to fusion. In *Proceedings of the 2Nd ACM SIGPLAN Workshop on Functional High-performance Computing*, FHPC '13, pages 47–58, New York, NY, USA, 2013. ACM.

[3] Hyperfit. `http://hiperfit.dk/`. Accessed: 05-15-2016.

[4] Troels Henriksen. Exploiting functional invariants to optimise parallelism: a dataflow approach. Master's thesis, DIKU, Denmark, 2014.